

# Pragmatic ways of using Rust in your data project

Christopher Prohm  
PyCon.DE / PyData Berlin 2023

[@chmp@hachyderm.io](mailto:@chmp@hachyderm.io)


Code + Slides: <https://github.com/chmp/PyConDE23>

# Motivation

## (+) Interactive exploration:

- data analysis
- model building
- writing tests
- ...


(-) Only fast if data fits NumPy, Pandas, ...  
& it's really easy to fall off the performance cliff



## Strategy: Do not replace Python

- Sprinkle Rust in for performance
- Rust & Python have complementary strengths

# How to turn raw data to “Python-compatible” data? (fast)



(\*) the Arrow format expands the range of what can be stored in data frame

[\[arrow.apache.org\]](https://arrow.apache.org)

## Examples:

1. Parsing bank statement PDFs
2. Converting the “Spotify Million Playlist Dataset” into a data frame

[\[aicrowd.com/challenges/spotify-million-playlist-dataset-challenge\]](https://aicrowd.com/challenges/spotify-million-playlist-dataset-challenge)

# Using Rust in your data project

**WARNING:** Rust Code ahead

# What I like about Rust

Performance & memory efficiency

Easy to integrate into other runtimes

- No garbage collector, no runtime
- Built to interface with C code

Well designed features fit together into high-level interfaces

- Macros (“Code generation”)
- Traits (“Interfaces”)
- Type inference
- Sum-types (“Unions”)
- ...

Great tooling & documentation!



# Python + Rust: PyO3 - the gold standard

PyO3 [\[pyo3.rs\]](https://pyo3.rs) allows to build Pythonic interfaces in Rust (\*)

Some Python libraries are fully written in Rust

Leverages Rust features to great effect:

- Macros for code generation
- Traits (“Interfaces”)
- Type inference

```
use pyo3::prelude::*;

#[pyfunction]
fn double(x: i32) -> i32 {
    x * 2
}

#[pymodule]
fn my_extension(py: Python<'_>, m: &PyModule) -> PyResult<()> {
    m.add_function(wrap_pyfunction!(double, m)?);
    Ok(())
}
```



Code: [pyo3.rs](https://pyo3.rs)

The hacky / easy alternative: building custom command line tools

(\*) See also [Rusty Python: A Case Study](#) by Robin Raymond

# Basic strategy (1): build custom CLI tools

CLI tool with JSON in, JSON out:

```
> echo '{"value": 21}' | io-patterns-double.exe  
{ "value": 42 }
```

JSON: lingua franca of data exchange

Potential alternative: binary encoding  
(e.g., bencode, Python code can be generated with  
serde-generate)

Using from Python

```
import subprocess  
  
res = subprocess.run(  
    ["/io-patterns-double.exe"],  
    encoding="utf-8",  
    capture_output=True, check=True,  
    input=inp,  
)  
out = res.stdout
```

Strings in Rust are  
UTF-8 encoded



How to encode the input / decode the output?

# Basic strategy (2): types + codegen in Python

```
from dataclasses import dataclass
```

```
@dataclass  
class Input:  
    value: int
```

```
@dataclass  
class Output:  
    value: int
```



`cattrs`<sub>[[catt.rs](#)]</sub>: simple & configurable way to convert between Python types and JSON compatible types

Simple cases require no custom code

No standardized encoding for unions

```
inp = Input(value=21)  
inp = cattrs.unstructure(inp)  
inp = json.dumps(inp)
```



`unstructure()`

Python  
object

JSON compatible  
objects

`structure()`

```
out = json.loads(out)  
out = cattrs.structure(out, Output)  
out
```





# Basic strategy (3): types + codegen in Rust

```
use serde::{Serialize, Deserialize};
```

```
#[derive(Deserialize)]
```

```
struct Input {  
    value: i64,  
}
```

```
#[derive(Serialize)]
```

```
struct Output {  
    value: i64,  
}
```



```
let input: Input = serde_json::from_reader(std::io::stdin())?;
```

```
let output = Output {  
    value: 2 * input.value,  
};
```

```
serde_json::to_writer(std::io::stdout(), &output)?;
```



Serde [[serde.rs](https://crates.io/crates/serde)]: de-facto Rust standard for (de)serialization

Rely on code generation via macros

# PyO3 vs. CLI Tools

## PyO3 [\[pyo3.rs\]](https://pyo3.rs)

- + High level wrapper around the Python C-API
  - rich conversions between Python & Rust
- + Maximum efficiency
  - shared objects
  - minimal call overhead
- Requires build & installation step
- No reloading (Python limitation)

## Custom CLI tool

- + Easy to get started:
  - No Rust / PyO3 specific concepts to learn
  - Easy to debug
- + Easy to build & distribute
- Not super efficient
- No shared objects

Code can still be refactored into PyO3

# Case study: Parsing PDFs

# Performance of PDF parsing

## Girokonto

Buchungstag Valuata	Auftraggeber / Empfänger IBAN/BIC	Ausgang Eingang
Alter Saldo		+XXXX,00
08.02.2021	Restaurant XYZ	-40,00
08.02.2021		
10.03.2021	Supermarkt XYZ	-30,00
10.03.2021		
...		
Neuer Saldo		+XXXX,00

PDF is a sequence of commands

```
Td 1.0 2.0
Tj "Alter Saldo"

Td 1.0 3.0
Tj "08.02.2021"
```

Needs to be interpreted

```
def do_Td(self, tx: PDFStackT, ty: PDFStackT):
    """Move text position"""

def do_Tj(self, s: PDFStackT):
    """Show text"""

...
```

Code from pdfminer.six

# Extracting structured information

Girokonto		
Buchungstag Valuata	Auftraggeber / Empfänger IBAN/BIC	Ausgang Eingang
Alter Saldo		+XXXX,00
08.02.2021	Restaurant XYZ	-40,00
08.02.2021		
.....		
10.03.2021	Supermarkt XYZ	-30,00
10.03.2021		
.....		
...		
Neuer Saldo		+XXXX,00

```
(1.0, 2.0, "Alter Saldo")
(1.0, 3.0, "08.02.2021")
(1.0, 4.0, "08.02.2021")
(12.0, 3.0, "-40,00")
(5.0, 3.0, "Restaurant XYZ")
...
```

```
[
    date(2021, 2, 8), -40.0),
    date(2021, 3, 10), -30.0),
    ...,
]
```



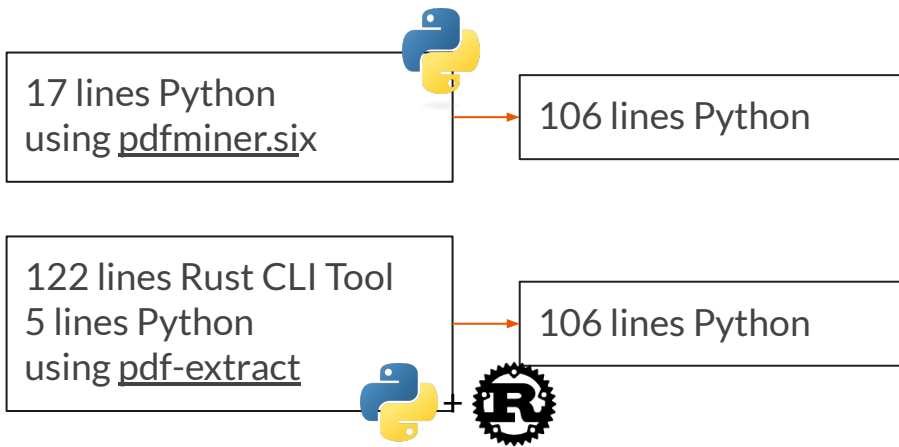
**Step 1:**  
Extract position  
+ text from PDF



**Step 2:**

1. Find header / footer of transactions
2. Group blocks
3. Parse dates, amounts

# Implementation



```
> pdf-parser.exe statement_2023-03-01.pdf
{"number":1,
 "words": [
   {"x": 1.0, "y": 2.0, "text": "..."},
   ...
 ]
}
```

For one year of statements:

25.1s

2.3s

11 x faster end to end

pdf-extract less mature than pdfminer.six

- e.g., content of the first page is not parsed

speed up depends on PDF: more complex -> bigger effect

# Case study: processing JSON files

# The Arrow revolution (\*) [\[arrow.apache.org\]](https://arrow.apache.org/)

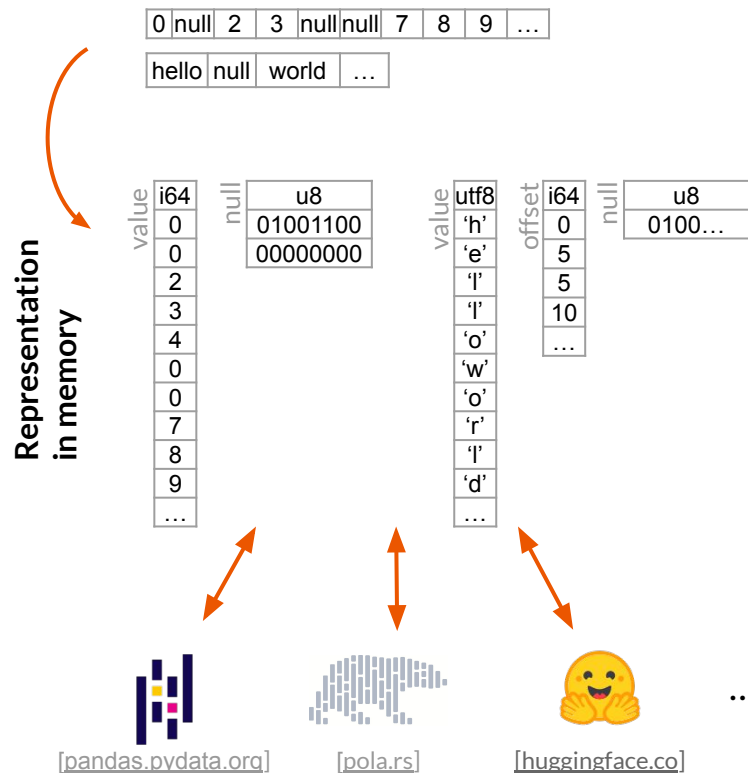
Specification how to arrange data frames in memory

Allows to exchange data without copies

Official support for C, C++, C#, Go, Java, JavaScript, Julia, MATLAB, **Python**, R, Ruby, **Rust**

Supports complex array types:

- Primitives (int64, float32, utf8, ...)
- Structs
- Lists
- ...



(\*) See also [Apache Arrow: connecting and accelerating dataframe libraries across the PyData ecosystem](#) by Joris Van den Bossche



# Processing JSON documents



Spotify Million Playlist Dataset

[\[aicrowd.com/challenges/spotify-million-playlist-dataset-challenge/\]](https://aicrowd.com/challenges/spotify-million-playlist-dataset-challenge/)

2.3 GB as Arrow IPC file


JSON documents with Spotify playlists

5.4 GB compressed, 30+ GB uncompressed

# Python Implementation

```
import pyarrow as pa


schema = {
    "name": pa.string(),
    'collaborative': pa.string(),
    'pid': pa.int64(),
    # ...
    'tracks': pa.list_(
        pa.struct({
            "pos": pa.int16(),
            'artist_name': pa.string(),
            'track_uri': pa.string(),
            # ...
        }),
    ),
}
```



```
with zipfile.ZipFile(root / "[...].zip", "r") as z:
    for i in range(1000):
        with z.open("...", "r") as fobj:
            d = json.load(fobj)

        for pl in d["playlists"]:
            pl["modified_at"] = 1000 * pl["modified_at"]

    table = pa.table({
        name: pa.array(
            [pl[name] for pl in d["playlists"]],
            type=ty,
        )
        for name, ty in schema.items()
    })
```



# Rust Implementation

```
#[derive(Deserialize, Serialize)]
```

```
struct Playlist {  
    name: String,  
    collaborative: String,  
    pid: i64,  
    modified_at: i64,  
    num_tracks: u16,  
    num_albums: u16,  
    num_followers: i64,  
    tracks: Vec<Track>,  
}
```



```
let mut builder = ArraysBuilder::new(&fields)?;
```

```
for i in 0..n {
```

```
    let mut content = Vec::new();
```

```
    zip.by_name("...",)? .read_to_end(&mut content)?;
```

```
    let mut data: Container = serde_json::from_slice(&content)?;
```

```
    for item in data.playlists.iter_mut() {  
        item.modified_at = 1000 * item.modified_at;  
    }
```

```
    builder.extend(&data.playlists)?;
```


```
}
```



Rust -> Arrow conversion: `serde_arrow` [[github.com/chmp/serde\\_arrow](https://github.com/chmp/serde_arrow)]

- Serde used in JSON and Arrow conversion
- Easy to write, but performance overhead
- Arrow 37.0.0 has similar feature built in

# Comparison



# Loading the data frame in Python using Polars [\[https://pola.rs\]](https://pola.rs)

```
import polars as pl
df = pl.read_ipc("data/spotify_million_playlist_dataset.ipc", memory_map=True)
```



```
(
    df.lazy().select(
        pl.col("tracks").arr.explode()
        .struct.field("artist_name")
        .value_counts().alias("counts")
    )
    .unnest("counts")
    .sort("counts").tail(10)
    .collect()
)
```



artist_name	counts
cat	u32
"J. Cole"	241560
"Justin Bieber"	243119
"Future"	250734
"Ed Sheeran"	272116
"Eminem"	294667
"The Weeknd"	316603
"Rihanna"	339570
"Kendrick Lamar..."	353624
"Kanye West"	413297
"Drake"	847160

# Extending Pandas / Polars / ...

## Built In PyO3 support in arrow-rs

```
#[pyfunction]
fn transform(array: &PyAny, py: Python) -> PyResult<PyObject> {
    let array = make_array(ArrayData::from_pyarrow(array)?);
    let result = todo!();
    result.to_data().to_pyarrow(py)
}
```

Converts only metadata  
Array data is shared



See also

[github.com/apache/arrow-rs/tree/master/arrow-pyarrow-integration-testing](https://github.com/apache/arrow-rs/tree/master/arrow-pyarrow-integration-testing)  
[docs.rs/arrow/latest/arrow/pyarrow/index.html](https://docs.rs/arrow/latest/arrow/pyarrow/index.html)

# Conclusion

# Conclusion

## Rust for data processing

- Rust is fast and memory efficient out of the box (no performance cliffs)
- Caveat: Rust libraries for data processing not yet as high quality as in Python

## Strategies

- Build your own CLI tools
- Use types & code generation to define interfaces
- Leverage Arrow to exchange structured data in a unified format
- Build extension modules with PyO3

## When to incorporate Rust?

- Performance benefit not always clear cut
- Bring data into a Python compatible format
- Complex processing steps (in particular string processing)
- Data not in data frame format

Code + Slides:  
[github.com/chmp/PyConDE23](https://github.com/chmp/PyConDE23)



# References


Code to talk: <https://github.com/chmp/PyConDE23>

Useful Rust libraries:

- [serde](#) - serialization & deserialization
- [serde json](#) - serialize into JSON
- [arrow](#) & [arrow2](#) - create arrow compatible data & write parquet
- [polars](#) - dataframes in Rust & Python
- [PyO3](#) - Python modules written in Rust
- [anyhow](#) - simplified error handling
- [rayon](#) - simple parallelization

# Backup

# Why Rust & Python?



Python builds on C, C++, Fortran

Rust is modern language, designed to fit into this group

# Streaming inputs / outputs (1)

```
> io-patterns-echo.exe
```

```
"foo"
```

```
"Echo: foo"
```

```
"bar"
```

```
"Echo: bar"
```

```
"baz"
```

```
"Echo: baz"
```

Often the output can be generated in parts

Strategy: use JSON lines / one line per message

Simple parallelization: Rust & Python can work in parallel

# Streaming inputs / outputs

```
with subprocess.Popen(  
    [path], encoding="utf-8",  
    stdin=subprocess.PIPE, stdout=subprocess.PIPE,  
) as proc:  
    # ...  
    proc.stdin.write(inp)  
    proc.stdin.write("\n")  
    proc.stdin.flush()  
  
    out = proc.stdout.readline()  
    # ...  
    proc.stdin.close() Signal end  
  
assert proc.returncode == 0
```

Deadlocks without  
flushing

```
std::io::stdout().write_all(&out)?;  
std::io::stdout().write_all(b"\n")?;  
std::io::stdout().flush()?;
```



Preventing deadlocks requires correct flushing &  
handling of EOF (end of file)




# Borrow checker: xor mutability

either-or:

- A **single** mutable reference
- Multiple immutable references

```
let mut s = String::from("hello");  
  
let r1 = &mut s;  
let r2 = &mut s;  
  
println!("{}", {}, r1, r2);
```

*won't compile*



Code: [doc.rust-lang.org](https://doc.rust-lang.org)

Tips:

- Consider cloning data if it simplifies the program
- Localize mutable access
- Split structs into smaller pieces



[github.com/luser/keep-calm-and-call-clone](https://github.com/luser/keep-calm-and-call-clone)

# Rust tip: keep it simple (1) (or don't to write Python in Rust)

```
struct S {  
    a: i32,  
    b: i32  
}  
  
impl S {  
    pub fn get_a(&mut self) -> &mut i32 {  
        &mut self.a  
    }  
    /* ... */  
}  
  
let mut s = S { a: 0, b: 0 };
```

```
fn update(a: &mut i32, b: &mut i32) {  
    *a = 1;  
    *b = 2;  
}
```

```
// compiles  
update(&mut s.a, &mut s.b);
```

```
// does not compile  
update(&mut s.get_a(), &mut s.get_b());
```

See also: [steveklabnik.com/writing/rusts-golden-rule](https://steveklabnik.com/writing/rusts-golden-rule)

# Rust tip: keep it simple (2) (or don't to write Python in Rust)

Python makes stream processing simple

```
def process_items(items: Iterable[T]) -> Iterable[U]:  
    for item in items:  
        ...  
        yield result
```



In Rust use out arguments

```
fn process_items(items: &[T], result: &mut Vec<U>) {  
    for item in items {  
        // ...  
        result.push(result)  
    }  
}
```



`&[T]` read-only view of a “list” of T

`&mut Vec<U>` write access to a “list” of U

Iterators can be written in Rust, but Lifetimes can make it complicated



# Rust tip: parallelize in Python

- Rust parallelism requirements are encoded in the type system (Send / Sync)
  - Often: independent units of works (e.g., files)
- ⇒ parallelize with Python threads

```
from concurrent.futures import ThreadPoolExecutor
from multiprocessing import cpu_count

with ThreadPoolExecutor(
    max_workers=cpu_count(),
) as executor:
    results = executor.map(process_files_in_rust, files)
```

